


# PROJECT 1

April 23, 2024



## 1 Business Understanding

Insurance companies invest a lot of time and money into optimizing their pricing and accurately estimating the likelihood that customers will make a [claim](#). In many countries insurance it is a legal requirement to have car insurance in order to drive a vehicle on public roads, so the market is very large!

Knowing all of this, On the Road car insurance have requested your services in building a model to predict whether a customer will make a claim on their insurance during the policy period. As they have very little expertise and infrastructure for deploying and monitoring machine learning models, they've asked you to identify the best performing model, as measured by accuracy, so they can start with the model in production.

They have supplied you with their customer data as a csv file called `car_insurance.csv`, along with a table detailing the column names and descriptions below.

### 1.1 The dataset

Column	Description
id	Unique client identifier
age	Client's age:
gender	Client's gender:
driving_experience	Years the client has been driving:
education	Client's level of education:
income	Client's income level:
credit_score	Client's credit score (between zero and one)
vehicle_ownership	Client's vehicle ownership status:
vehcile_year	Year of vehicle registration:
married	Client's marital status:
children	Client's number of children
postal_code	Client's postal code
annual_mileage	Number of miles driven by the client each year
vehicle_type	Type of car:
speeding_violations	Total number of speeding violations received by the client
duis	Number of times the client has been caught driving under the influence of alcohol
past_accidents	Total number of previous accidents the client has been involved in
outcome	Whether the client made a claim on their car insurance (response variable):

```
[105]: # Modules to handle table-like data and matrices
import pandas as pd
import numpy as np

# Modelling Algorithms Modules
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression

# Modelling Helpers Modules
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import LabelEncoder

# Visualization Modules
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[106]: cars = pd.read_csv('car_insurance.csv')
cars.head()
```

```
[106]:      id  age  gender  driving_experience  education  income \
0  569520   3      0             0-9y  high school  upper class
```

1	750365	0	1	0-9y	none	poverty
2	199901	0	0	0-9y	high school	working class
3	478866	0	1	0-9y	university	working class
4	731664	1	1	10-19y	none	working class

	credit_score	vehicle_ownership	vehicle_year	married	children	\
0	0.629027		1.0 after 2015	0.0	1.0	
1	0.357757		0.0 before 2015	0.0	0.0	
2	0.493146		1.0 before 2015	0.0	0.0	
3	0.206013		1.0 before 2015	0.0	1.0	
4	0.388366		1.0 before 2015	0.0	0.0	

	postal_code	annual_mileage	vehicle_type	speeding_violations	duis	\
0	10238	12000.0	sedan		0	0
1	10238	16000.0	sedan		0	0
2	10238	11000.0	sedan		0	0
3	32765	11000.0	sedan		0	0
4	32765	12000.0	sedan		2	0

	past_accidents	outcome
0	0	0.0
1	0	1.0
2	0	0.0
3	0	0.0
4	1	1.0

```
[107]: cars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    10000 non-null  int64
1   age                   10000 non-null  int64
2   gender                10000 non-null  int64
3   driving_experience    10000 non-null  object
4   education             10000 non-null  object
5   income               10000 non-null  object
6   credit_score          9018 non-null   float64
7   vehicle_ownership    10000 non-null  float64
8   vehicle_year         10000 non-null  object
9   married              10000 non-null  float64
10  children              10000 non-null  float64
11  postal_code           10000 non-null  int64
12  annual_mileage        9043 non-null   float64
13  vehicle_type          10000 non-null  object
14  speeding_violations   10000 non-null  int64
```

```
15  dui      10000 non-null  int64
16  past_accidents  10000 non-null  int64
17  outcome   10000 non-null  float64
dtypes: float64(6), int64(7), object(5)
memory usage: 1.4+ MB
```


## 2 Data Understanding

Check the balance between outcome

1 : claims made, 0 : no claims

```
[108]: sns.countplot(x='outcome', data=cars)
```


```
[108]: <Axes: xlabel='outcome', ylabel='count'>
```



Check the presence and distribution of outliers using boxplot


```
[109]: sns.boxplot(x='outcome', y='annual_mileage', data=cars)
```

```
[109]: <Axes: xlabel='outcome', ylabel='annual_mileage'>
```




```
[110]: sns.histplot(cars['credit_score'], bins=20, kde=True)
```

```
[110]: <Axes: xlabel='credit_score', ylabel='Count'>
```



```
[111]: sns.histplot(cars['annual_mileage'], bins=20, kde=True)
```

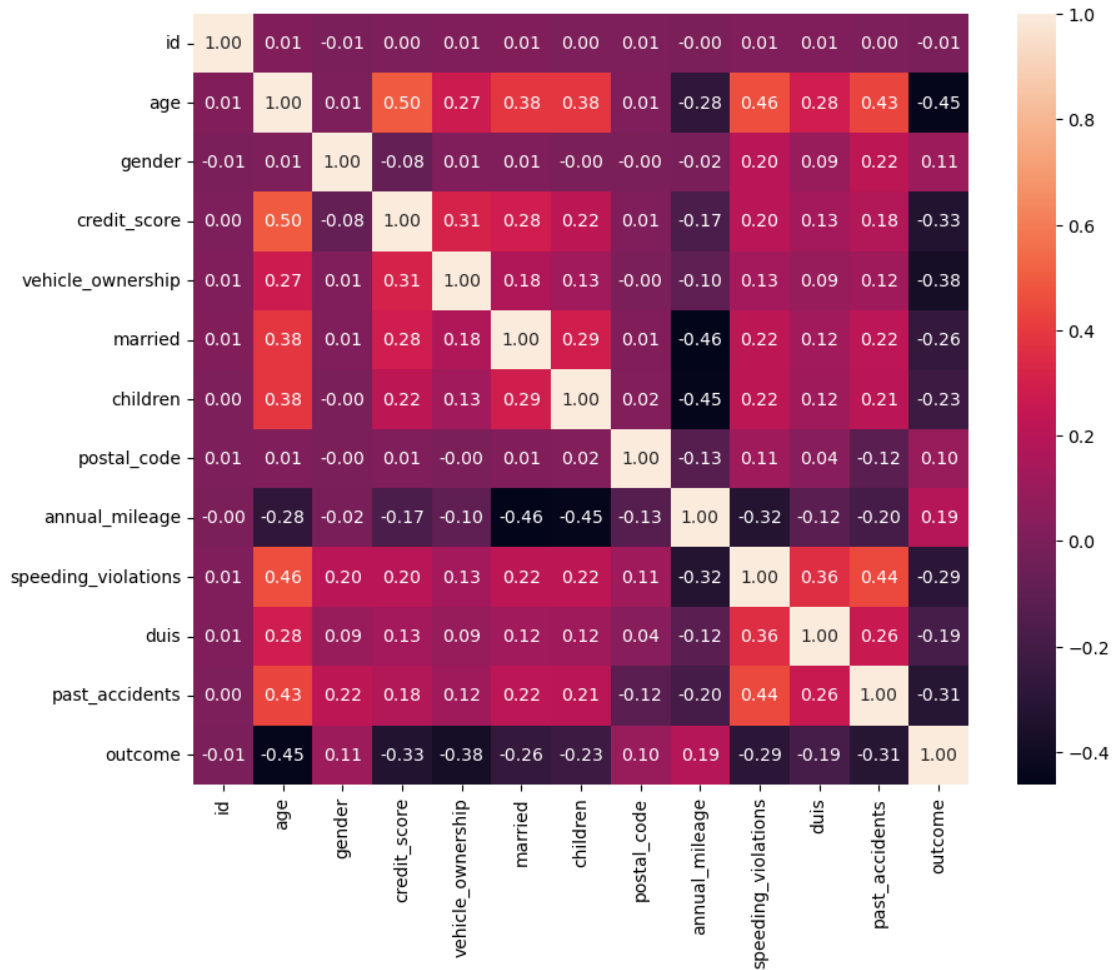
```
[111]: <Axes: xlabel='annual_mileage', ylabel='Count'>
```



```
[112]: correlation = cars.corr()
plt.figure(figsize=(10,8))
sns.heatmap(correlation, annot=True, fmt='.2f')
```

C:\Users\Aryo Sasi\AppData\Local\Temp\ipykernel\_2436\2857633541.py:1:  
FutureWarning: The default value of numeric\_only in DataFrame.corr is  
deprecated. In a future version, it will default to False. Select only valid  
columns or specify the value of numeric\_only to silence this warning.  
correlation = cars.corr()

```
[112]: <Axes: >
```



### 3 Data Preparation

handle missing value with mean function

```
[113]: cars['credit_score'].fillna(cars['credit_score'].mean(), inplace = True)
cars['annual_mileage'].fillna(cars['annual_mileage'].mean(), inplace = True)
cars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    10000 non-null  int64
1   age                   10000 non-null  int64
2   gender                10000 non-null  int64
3   driving_experience     10000 non-null  object
```



```

4  education          10000 non-null object
5  income             10000 non-null object
6  credit_score       10000 non-null float64
7  vehicle_ownership  10000 non-null float64
8  vehicle_year       10000 non-null object
9  married            10000 non-null float64
10 children           10000 non-null float64
11 postal_code        10000 non-null int64
12 annual_mileage     10000 non-null float64
13 vehicle_type       10000 non-null object
14 speeding_violations 10000 non-null int64
15 duis               10000 non-null int64
16 past_accidents     10000 non-null int64
17 outcome            10000 non-null float64
dtypes: float64(6), int64(7), object(5)
memory usage: 1.4+ MB

```

Handle outliers in the annual\_mileage column using RobustScaler

```

[114]: from sklearn.preprocessing import RobustScaler

scaler = RobustScaler()
cars['annual_mileage_scaled'] = scaler.fit_transform(cars[['annual_mileage']])
cars.head()

```

```

[114]:      id  age  gender  driving_experience  education  income \
0  569520    3      0           0-9y  high school  upper class
1  750365    0      1           0-9y      none  poverty
2  199901    0      0           0-9y  high school  working class
3  478866    0      1           0-9y  university  working class
4  731664    1      1          10-19y      none  working class

      credit_score  vehicle_ownership  vehicle_year  married  children \
0      0.629027          1.0  after 2015      0.0      1.0
1      0.357757          0.0  before 2015      0.0      0.0
2      0.493146          1.0  before 2015      0.0      0.0
3      0.206013          1.0  before 2015      0.0      1.0
4      0.388366          1.0  before 2015      0.0      0.0

      postal_code  annual_mileage  vehicle_type  speeding_violations  duis \
0      10238          12000.0      sedan          0      0
1      10238          16000.0      sedan          0      0
2      10238          11000.0      sedan          0      0
3      32765          11000.0      sedan          0      0
4      32765          12000.0      sedan          2      0

      past_accidents  outcome  annual_mileage_scaled
0              0      0.0          0.100999

```

1	0	1.0	1.434332
2	0	0.0	-0.232334
3	0	0.0	-0.232334
4	1	1.0	0.100999

Create a function to encode categorical columns using LabelEncoder

```
[115]: def columns_le(data, columns):
        le = LabelEncoder()
        for column in columns:
            data[f'{column}_encoded'] = le.fit_transform(data[column]) # encode the
            ↪column
            data.drop(column, axis = 1, inplace=True) # drop the previous column
        return data
```

```
[116]: columns_to_encode = ['driving_experience', 'education', 'income',
            ↪'vehicle_year', 'vehicle_type']
        encoded_cars = columns_le(cars, columns_to_encode)
        encoded_cars.drop('annual_mileage', axis=1, inplace=True)
        encoded_cars.head()
```

```
[116]:      id  age  gender  credit_score  vehicle_ownership  married  children \
0  569520    3      0      0.629027              1.0         0.0         1.0
1  750365    0      1      0.357757              0.0         0.0         0.0
2  199901    0      0      0.493146              1.0         0.0         0.0
3  478866    0      1      0.206013              1.0         0.0         1.0
4  731664    1      1      0.388366              1.0         0.0         0.0
```

	postal_code	speeding_violations	duis	past_accidents	outcome	\
0	10238	0	0	0	0.0	
1	10238	0	0	0	1.0	
2	10238	0	0	0	0.0	
3	32765	0	0	0	0.0	
4	32765	2	0	1	1.0	

	annual_mileage_scaled	driving_experience_encoded	education_encoded	\
0	0.100999	0	0	
1	1.434332	0	1	
2	-0.232334	0	0	
3	-0.232334	0	2	
4	0.100999	1	1	

	income_encoded	vehicle_year_encoded	vehicle_type_encoded
0	2	0	0
1	1	1	0
2	3	1	0
3	3	1	0



Performing Hyperparameter tuning using GridSearch to find the best paramaters and model.

```
[119]: model_params = {
        'svc': {
            'model': SVC(gamma='auto'),
            'params' : {
                'C': [1,10,20]
            }
        },
        'random_forest': {
            'model': RandomForestClassifier(),
            'params' : {
                'n_estimators': [1,5,10]
            }
        },
        'logistic_regression' : {
            'model': LogisticRegression(solver='liblinear',multi_class='auto'),
            'params': {
                'C': [1,5,10]
            }
        }
    }
```

## 5 Evaluation

```
[127]: scores = []

for model_name, mp in model_params.items():
    clf = GridSearchCV(mp['model'], mp['params'], cv = 5,
    ↪return_train_score=False)
    clf.fit(X_train, y_train)
    scores.append({
        'model':model_name,
        'best_score':clf.best_score_,
        'best_params':clf.best_params_
    })
```

```
[129]: df = pd.DataFrame(scores, columns=['model', 'best_score', 'best_params'])
df
```

```
[129]:
```

	model	best_score	best_params \
0	svc	0.848000	{'C': 1}
1	random_forest	0.823125	{'n_estimators': 10}
2	logistic_regression	0.820625	{'C': 1}

	best_estimator	best_index
0	SVC(C=1, gamma='auto')	0

```

1 (DecisionTreeClassifier(max_features='sqrt', r...      2
2     LogisticRegression(C=1, solver='liblinear')      0

```

Build a simple prediction system

using the first row index where id is 569520 and the outcome is 0 : no claims

```
[122]: encoded_cars.head()
```

```
[122]:
```

	id	age	gender	credit_score	vehicle_ownership	married	children	\
0	569520	3	0	0.629027	1.0	0.0	1.0	
1	750365	0	1	0.357757	0.0	0.0	0.0	
2	199901	0	0	0.493146	1.0	0.0	0.0	
3	478866	0	1	0.206013	1.0	0.0	1.0	
4	731664	1	1	0.388366	1.0	0.0	0.0	

	postal_code	speeding_violations	duis	past_accidents	outcome	\
0	10238	0	0	0	0.0	
1	10238	0	0	0	1.0	
2	10238	0	0	0	0.0	
3	32765	0	0	0	0.0	
4	32765	2	0	1	1.0	

	annual_mileage_scaled	driving_experience_encoded	education_encoded	\
0	0.100999	0	0	
1	1.434332	0	1	
2	-0.232334	0	0	
3	-0.232334	0	2	
4	0.100999	1	1	

	income_encoded	vehicle_year_encoded	vehicle_type_encoded
0	2	0	0
1	1	1	0
2	3	1	0
3	3	1	0
4	3	1	0

```
[123]: input_data = (3, 0, 0.629027, 1.0, 0.0, 1.0, 10238, 12000.0, 0, 0, 0, 0, 0, 2, 0, 0)
```

```
svc = SVC(C=1, gamma='auto')
```

```
input_array = np.asarray(input_data)
```

```
input_resaped = input_array.reshape(1,-1)
```

```
svc.fit(X_train, y_train)
```

```
prediction = svc.predict(input_resaped)
print(prediction)
```

[0.]

```
C:\Users\Aryo Sasi\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning:
X does not have valid feature names, but SVC was fitted with feature names
warnings.warn(
```

## 6 Conclusion

The purpose of this project is to build a model that predicts whether a customer will make an insurance claim. After testing several models, the best-performing one was the Support Vector Machine with a parameter C of 1, which achieved the highest accuracy of 84% in predicting customer insurance claim submissions. In testing simple predictions, the model also correctly predicted No Claims for the first row.



[ ]: